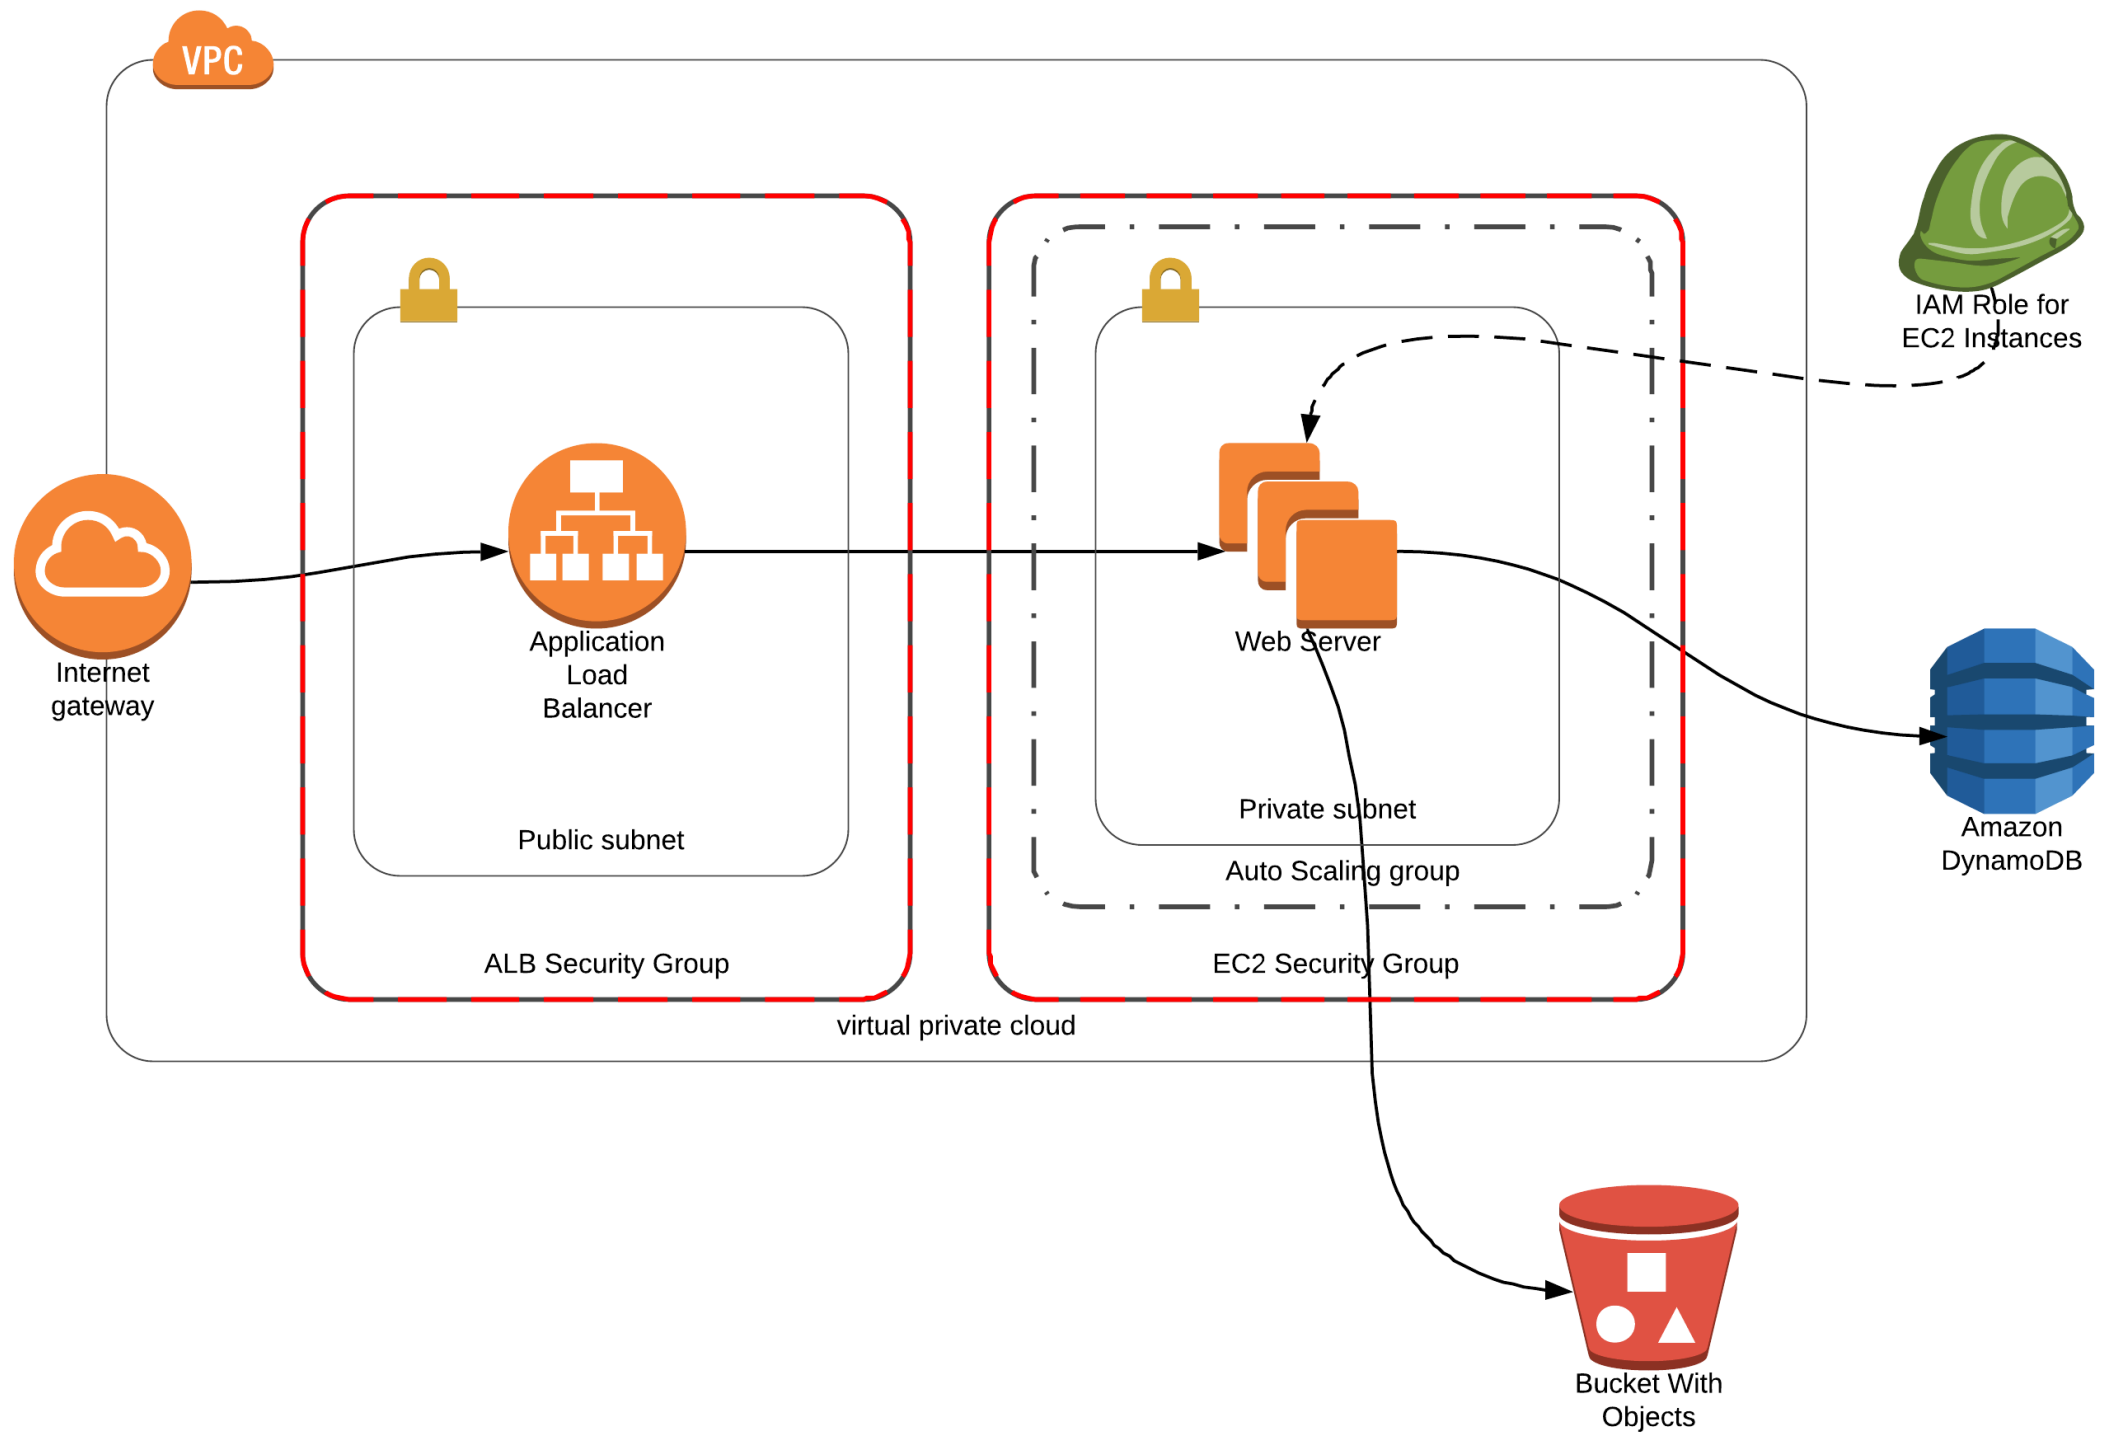


Introduction to Terraform and Terraform Modules

Alberto Castro

What is Terraform?

- “Terraform is a tool for **building, changing, and versioning infrastructure safely and efficiently.**”
- **Building, changing, and versioning** – *Infrastructure as Code that can be managed with Git.*
- **Safely and efficiently** – Terraform maintains state of infrastructure and applies the changes in a (mostly) non-conflicting manner.



Goal: Build this infrastructure with Terraform

Terraform Configuration for AWS

- Install Terraform downloading it from <https://www.terraform.io/downloads.html>
- Install `awscli` in your computer.
- To configure terraform create a user in AWS using IAM and name it `terraform` and programatic access (no console login)
- Assign full administration permissions to the user.
- Download the Access Keys and upload them and initialize the awscli by running:
`aws configure`

Terraform Basic Use

- Create a terraform file in a directory in your computer.
- If it is the first time you are running terraform for this script you should run `terraform init` so that terraform can download the corresponding providers (aws in our case)
- Run `terraform plan` and check the changes that will be applied.
- Run `terraform apply` and after accepting the changes terraform will provision your infrastructure.
- If you desire to destroy the infrastructure run `terraform destroy` this will remove all the resources created before.

Other Useful Commands

- `terraform fmt` Formats the terraform files.
- `terraform state subcommand` Allow you to display and manage terraform state (`list`, `rm`, `show`, ...)
- Run `terraform validate` Validates the Terraform files.
- `terraform taint` Manually mark a resource for recreation
- `terraform untaint` Manually unmark a resource for recreation

Example 1 - Basic

Homework

- Create an AWS account, create a terraform user and generate a key pair for it.
- Initialize terraform and deploy example1 draft1, draft2 and final. Examine the result in the console. Remember to run terraform destroy when you finish working to avoid extra charges.
- Using example1/draft1 deploy one more ec2 instance by copying the definition and configuring it differently.


- Using the `aws_security_group_rule` resource rewrite the ingress rules for the aws security group `ec2_security_group` defined in main (remember to remove the `ingress` rules in the `aws_security_group_rule` definition).
- Re-run terraform with this configuration. The result should be:
 - If you already deployed the terraform resources terraform plan should tell you that no changes are necessary.
 - If you are deploying them anew you should get the same result as the final example.

Terraform Architecture

- Terraform is written in **Go** and it is a tool developed and maintained by Hashicorp (<https://www.hashicorp.com/>).
- A **core engine** manages the language features, so that the same language (HCL) is used for all infrastructure types.
- **Providers** create the resources in the different cloud services. This way depending on the provider loaded different infrastructure can be managed.
- The AWS Terraform provider leverages the **Go SDK** to create resources.

Language

- Terraform's language is *declarative*. A terraform script specifies what is the final state of the infrastructure after changes are applied.
- Terraform manages all dependencies (mostly), creating resources in the correct order.
- The latest version of Terraform (0.12) uses HCL 2.0 which is a major improvement over the previous version making the code more readable and easier to write.
- Every script is its own “module” which in Terraform means a script that accepts inputs and outputs (like a function) and that can be linked to other modules.



Terraform Module Structure

- Iteration can be achieved using a **counter** and (more recently) **for loops**.
- There is a type structure for variables with:
 - Primitive Types: **string**, **number**, **bool**
 - Collection Types: **list**, **map**, **set** (all of same type elements)
 - Structural Types: **tuple** (a lists of multiple types) and **object** (a map of multiple types)
- All these types have literal expressions to declare explicit values.

- Primitive type examples: `"this is a string"`; `12345` (my favorite password!); `true`
- `List`: `["a", "b", "z"]`
- `Map`:

```
{  
  us-east-1    = "ami-0a01a5636f3c4f21c"  
  us-east-2    = "ami-0a02eadc6d8770f83"  
  us-west-1    = "ami-0bbeea654a35ef611"  
  us-west-2    = "ami-0a1af68029fa293b6"  
}
```
- `Object`:

```
{  
  anumber = 12  
  alist   = ["2", "w", "foo"]  
  atuple  = [2, false, "baz"]  
}
```

Example 2 - Modules

Terraform State

- Terraform maintains a file (terraform.tfstate) that contains the state of the infrastructure after the last applied change.
- Initially Terraform assumes that the resources that it is creating do not exist, and as new resources are created the terraform.tfstate file is updated.
- It is possible to import state, although it is cumbersome.
- It is *highly discouraged* to make changes without Terraform once it has been created with it. Doing so means having to import the new state or give up on using Terraform altogether.

- Terraform State stores information about the current resources, the parameters, and dependencies.
- API calls are made to the cloud provider to check the infrastructure state and compare it with the changes.
- The state *can* be used by terraform to plan if desired, this is useful for very large infrastructure deployments.
- The state can be saved and shared using for example:
 - Terraform Cloud
 - S3 bucket + DynamoDB table for locking

Storing and Sharing Modules:

1. The Registry

- Terraform registry is a repository of modules freely available for use directly from terraform.
- All modules are versioned and fully documented.
- Modules are available for Terraform v0.11 and v0.12. Examples are for Terraform v0.12 and HCL 2.0 in this presentation.
- Some modules are “verified” which means that the modules have been written and reviewed by Hashicorp engineers.

Storing and Sharing Modules:

2. Git Repositories

- Git repositories and in particular GitHub can be used to store and share repositories. (<https://www.terraform.io/docs/modules/sources.html>)
- Modules defined in a git repository are loaded with:
`source = "git::https://example.com/vpc.git"`
- Branches and tags can be specified allowing for versioning and source control. Here `v1.2.4` is the name of a branch or tag:
`source = "git::https://example.com/vpc.git?ref=v1.2.4"`

Terraform in AWS

- Terraform is recommended by AWS on par with their own tool CloudFormation to manage and deploy infrastructure in AWS.
- Terraform's AWS provisioner receives frequent updates and new features are deployed shortly after, or sometimes as the services are added to AWS.
- It is written with the AWS Go SDK so as features are added there they are added to Terraform.
- Terraform is used by AWS engineers.

How to learn Terraform

- Learn first some Cloud Infrastructure principles (Virtual Servers, IAM, Databases, VPCs, etc).
- Open a free account in a major cloud provider and play with the console for a while.
- Recreate the infrastructure created before using Terraform.
- After that check out code in the official Terraform registry and study it.
- Find a cool project and deploy your infrastructure with Terraform, do not worry, you can start small and add more later.

Resources

- Github Repository with the examples and slides
<https://github.com/albertocastro63/terraform-nerd>
- Terraform Web Site
<https://www.terraform.io/>
- Terraform Registry
<https://registry.terraform.io/>
- *Terraform Up And Running 2nd Edition* by **Yevgeniy Brikman** O'Reilly Media, September 2019
<http://shop.oreilly.com/product/0636920225010.do>